

Mitigating Tony Hoare's CAD\$2.16 billion mistake

C++ Ottawa
27 March 2025
Marc Pawlowsky

2009



Null References: The Billion Dollar Mistake - Tony Hoare

Tony Hoare introduced Null references in ALGOL W back in 1965 "simply because it was so easy to implement"

<https://www.infoq.com/presentations/Null-References-The-Billion-Dollar-Mistake-Tony-Hoare/>

What is `int* p`?

- Pointer to a piece of data
- Pointer to a collection of data, `*(p+4)`
- No data is available,
`int *p = 0; p == nullptr;`
- Pointer to a resource that must be released
- Pointer to a resource that must not be released

```
void  
do_not_delete(int const * const p)  
{  
    delete p;  
}
```

<https://godbolt.org/z/zMd6YsqMr>

Today's goals

- Reduce the chance of following a null-pointer
- Reduce your cognitive load
 - Make the implicit explicit

Backlog

- Pointer to a piece of data
- Pointer to a collection of data, $*(p+4)$
- 0 as nullptr
- No data is available
- Pointer to a resource that must be released
- Pointer to a resource that must **not** be released

Backlog

- Pointer to a piece of data
- Pointer to a collection of data, $*(p+4)$
- 0 as nullptr
- No data is available
- Pointer to a resource that must be released
- Pointer to a resource that must **not** be released
- **What is the binary representation of 0**

Inspiration

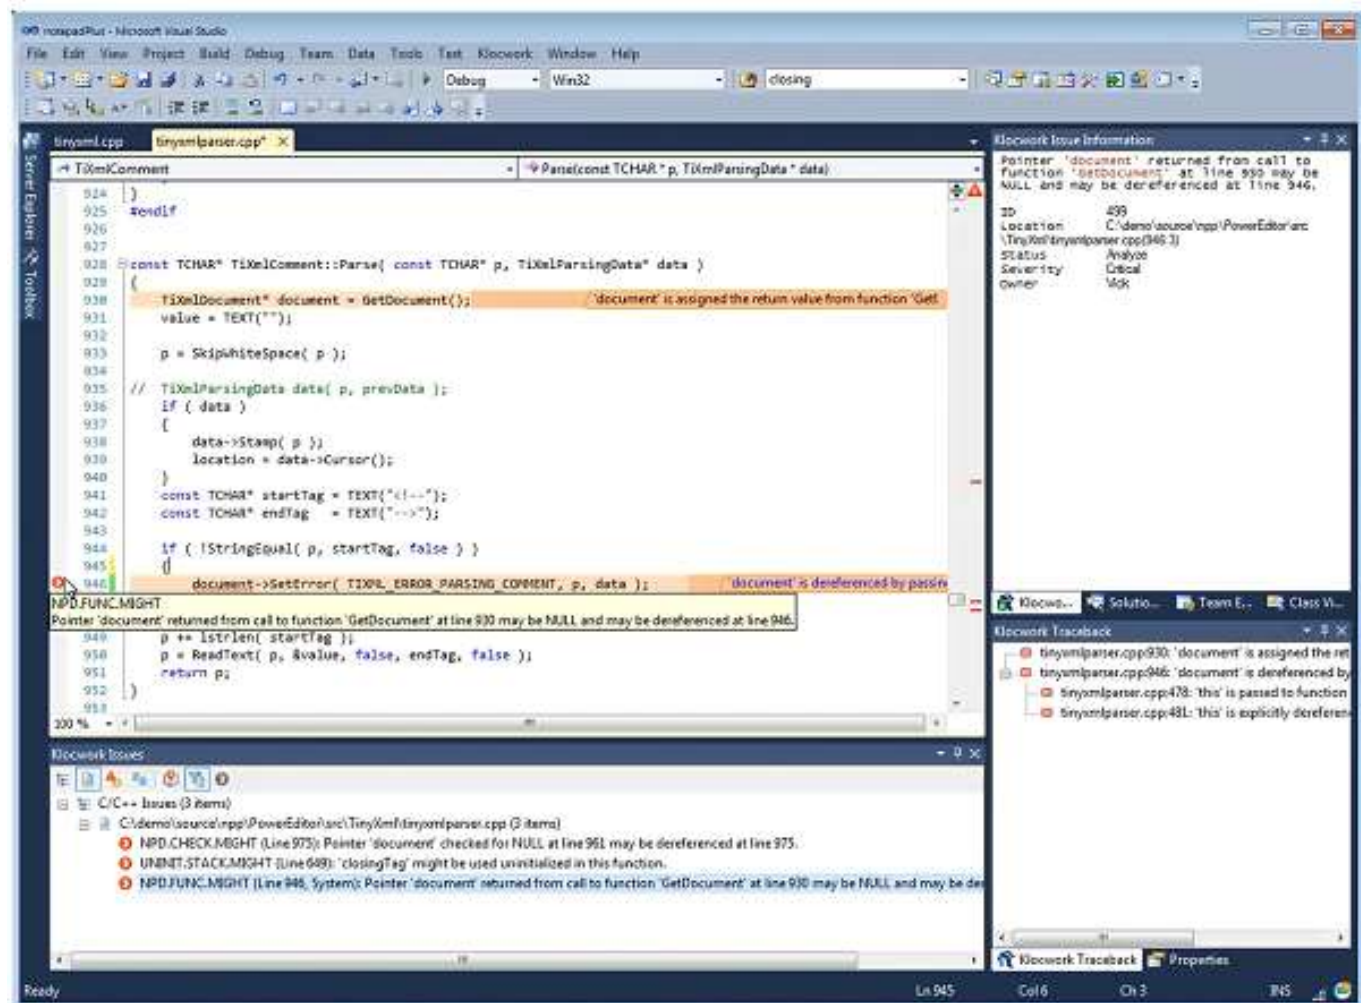
- Java Programmers FAQ post by Marc Pawlowsky (~1996)
Does Java have the equivalent of "const" arguments in C and C++?
- Correct by Construction: APIs That Are Easy to Use and Hard to Misuse :
C++ On Sea, Matt Godbolt, 2020
- [gsl::not_null](#), and [gsl::owner](#)
Guidelines support library

Technique


- Step by step refactoring of legacy code to catch the possible misuse pointers
 - Real world:
 - C++98 using Boost
 - New code in C++17
- Introduce a compiler failure in an API and fixit.

Tools?

- **GSL** : Guidelines Support Library
- clang-tidy
- static analyzers



Relative Cost of Fixing Defects



mp

- namespace mp = marcpawl::pointers
- <https://github.com/marcpawl/nullptr/tree/20250327>



Wrapping gotchas

- `operator bool()`

Take a look at:

2020 LLVM Developers' Meeting

Using Clang-tidy for Customized Checkers and Large Scale Source Tree Refactoring - Vince Bridgers

- How to write your own rules
- Fix one clang-tidy rule at a time across a whole project

Step 1: nullptr

- Replace 0 with nullptr
 - `clang-tidy - modernize-use-nullptr`

```
int* p = 0;  
=====>  
int* p = nullptr;
```


Step 1: nullptr

- Replace implicit use in conditional
 - **clang-tidy - readability-implicit-bool-conversion**

```
int* p;  
If (p) {}  
=====>  
if (p != nullptr) {}
```

Backlog

- Pointer to a piece of data
- Pointer to a collection of data, $*(p+4)$
- ~~0 as nullptr~~
- No data is available
- Pointer to a resource that must be released
- Pointer to a resource that must **not** be released
- What is the binary representation of 0

Step 2 : Describe data availability

Data may not available

Comments as a type

```
// a flag for saying go fetch the data
void foo(int* ptr) {
    if (nullptr == ptr) {
        ptr = new int(4);
    }
}
```

```
// a flag for no-op
void fee(int* ptr) {
    if (nullptr == ptr) {
        return;
    }
}
```

Data must be available

Comment as type

```
// @param p may not be null
int square(int* p) {
    return (*p) * (*p);
}
```

Data must be available

Soft failure

```
// @param p must not be null, ...
// @return -1 if p is nullptr, else ...
int square(int* p) {
    if (nullptr != p) {
        // Fix for BUG12345
        return -1; // in-band encoding
    }
    return (*p) * (*p);
}
```

Data must be available

Hard failure

```
// @param p must not be null
int square(int* p) {
    assert(nullptr != p);
    return (*p) * (*p);
}
```

Backlog

- ~~Pointer to a piece of data~~
- **Runtime type description of data must be available**
- Pointer to a collection of data, $*(p+4)$
- ~~No data is available~~
- Pointer to a resource that must be released
- Pointer to a resource that must **not** be released
- What is the binary representation of 0

Step 3: Use a type

- Have the compiler tell you when data must be available with run time checks.

Step 3.1 : gsl::not_null

- Use `gsl::not_null` to indicate data MUST be available

```
void hum(gsl::not_null<int*> p)
```

- Opinionated: prevents pointer arithmetic

- `clang-tidy - cppcoreguidelines-owning-memory`

<https://github.com/microsoft/GSL/blob/3.0.1/include/gsl/pointers#L69>

Step 3.2 : gsl::strict_not_null

- Use `gsl::strict_not_null`
- Removes explicit constructor and automatic conversion.

```
void hum(gsl::strict_not_null<int*> p)
```

Cost gsl::not_null

- No extra space
- ```
void foo(int *p) { auto np = gsl::strict_not_null<int*>(p);}
```

  
we have a **conditional**, equivalent to the assert.
- ```
void bar(gsl::strict_not_null<int*> q);  
void foo(gsl::strict_not_null<int*> p) { bar(p);}
```


No conditional, no need for asserts.
- Compiler **optimizes** out extra check.

Step 3.3 : mp::nullable

Type as a comment.

```
template<
    class T,
    std::enable_if_t<std::is_pointer<T>::value, bool> = true>
using nullable = T;

void foo(nullable<int*> ptr);
```

Step 3.4 Overload pointer functions

```
void foo(mp::maybe_null<int*> p)
{
    if (p == nullptr) {
        return;
    }
    // DO something big with p
}
```

```
void foo(gsl::not_null<int*> p)
{
    // DO something big with p
}
```

```
void foo(mp::maybe_null<int*> p)
{
    if (p == nullptr) {
        return;
    }
    foo(gsl::make_not_null(p));
}
```

Where are we

- `gsl::strict_not_null<int*>`
- `mp::nullable<int*>`
 - it is intentional that null is a valid value
- `int*`
 - To be determined

Backlog

- ~~Compile time description of data must be available~~
- **Force nullable to be checked before use.**
- Pointer to a collection of data, $*(p+4)$
- Pointer to a resource that must be released
- Pointer to a resource that must **not** be released
- What is the binary representation of 0

Step 4 : std::span or gsl::span

- [clang-tidy - cppcoreguidelines-pro-bounds-pointer-arithmetic](#)
- Switch from legacy:
`void print(Node const* nodes, size_t count)`

To bounds check:
`void print(std::span<Node const>& nodes)`

<https://godbolt.org/z/71Tn5s3KE>

- `std::span<T>::operator[](size_t)`
Bounds checking
Coming in C++26

std::span cost

- Space is larger?
 - Pointer
 - Size
 - Maybe you only need a UINT8

Backlog

- Force nullable to be checked before use.
- ~~Pointer to a collection of data, $*(p+4)$~~
- Pointer to a resource that must be released
- Pointer to a resource that must **not** be released
- What is the binary representation of 0

Delete or not to delete?

```
struct Node {  
    Node* next;  
    Node(Node* next_):  
        next(next_)  
    {}  
    ~Node() {  
        delete next; // ?  
    }  
}
```

Step 5 : Comments as a type

```
struct Node {  
    // next's lifespan is distinct from this node.  
    // should not be deleted.  
    Node* next;  
    Node(Node* next_):  
        next(next_)  
    {}  
    ~Node() {  
    }  
}
```

Backlog

- Force nullable to be checked before use.
- ~~Pointer to a resource that must be released~~
- ~~Pointer to a resource that must **not** be released~~
- **Type for Pointer to a resource that must be released**
- **Type for Pointer to a resource that must not be released**
- What is the binary representation of 0

Step 6 : types for ownership

Types as comments

- `gsl::owner`

```
template <class T, std::enable_if_t<std::is_pointer<T>::value,  
bool> = true>  
using owner = T;
```

<https://godbolt.org/z/ax45bfYxz>

- [gsl/pointers#L78](https://godbolt.org/z/ax45bfYxz)

Types as comments

- `mp::nonowner`

<https://godbolt.org/z/raeTMcWvj>

```
template <
  class T,
  std::enable_if_t<std::is_pointer<T>::value, bool> = true>
using nonowner = T;

void bar(nonowner<Node*> t) { ...}
```

Owner guidelines

- Do not use `gsl::owner<T*> const&`, use `mp::nonowner<T*> const&` instead.

If you are not manipulating the pointer why are you identifying it as an owner.

Where are we

<code>gsl::owner< gsl::strict_not_null<int*>></code>	<code>gsl::owner< mp::nullable<int*>></code>	<code>gsl::owner< std::span<int>></code>
<code>mp::borrower< gsl::strict_not_null<int*>></code>	<code>mp::nonowner< mp::nullable<int*>></code>	<code>mp::nonowner< std::span<int>></code>
<code>gsl::strict_not_null<int*></code>	<code>mp::nullable<int*></code>	<code>std::span<int></code>
<code>gsl::owner<int*></code>	<code>mp::nonowner<int*></code>	<code>int*</code>

Backlog

- Force nullable to be checked before use.
- ~~Type for Pointer to a resource that must be released~~
- ~~Type for Pointer to a resource that must not be released~~
- **Managed resources**
- What is the binary representation of 0

Step 7 : Managed resources

- `std::unique_ptr`
- `std::shared_ptr`
- `std::weak_ptr`

std::shared_ptr

- <https://godbolt.org/z/z1hnxKzxE>

std::weak_ptr

<https://godbolt.org/z/M1rczeMjd>

- Use lock() to get a shared_ptr
- Check to see if not null_ptr

std::shared_ptr

- Reference counted
 - Not garbage collected
- Thread safe initialization
 - Slow to construct
- Indicates shared ownership

std::unique_ptr

- Used to indicate ownership

<https://godbolt.org/z/Ecd1sx41P>

Non-null managed pointers

- `gsl::strict_not_null<std::unique_ptr<int>&> &`
- `gsl::strict_not_null<std::shared_ptr<int>> &`

<https://godbolt.org/z/Gs59Pzran>

nullable pointers

- `std::optional<gsl::strict_not_null<int*>> &`

<https://godbolt.org/z/1Y7roG44d>

nullable pointers

```
using VN = std::variant<gsl::strict_not_null<int*>, std::nullptr_t>;  
VN factory(gsl::strict_not_null<int*> b)
```

Use `std::visit` and you are guaranteed to handle both cases


<https://godbolt.org/z/oxxTMzf71>

Backlog

- ~~Force nullable to be checked before use.~~
- ~~Managed resources~~
- **Cheaper checks on nullable**
- **Enforce safety owner = borrower**
- What is the binary representation of 0

One step beyond

Madness



Step 8: Compile time checking

Haskell

```
data OptionalInt = Some Int | None  
    deriving (Show)
```

```
addOne :: OptionalInt -> OptionalInt  
addOne None      = None  
addOne (Some x) = Some (x + 1)
```

mp::strict_not_null

- copy gsl::strict_not_null
- Extra overloads to handle
std::optional<strict_not_null<std::unique_ptr>>

mp::maybe_null

- Deprecates *, ->
- as_optional_not_null
- as_variant_not_null
- visit

[ptr.hpp#L488](#)

[maybe_null_tests.cpp#L358](#)

mp::borrower and mp::owner

- Owner cannot share ownership
- Owner can share usage with borrower
- Borrower can share usage with another borrower
- Borrower cannot give away ownership

mp::owner

- Cannot be assigned to another owner

```
auto fail()
{
    int p;
    mp::borrower<int*> ptr{&p};
    // ERROR: cannot create an owner from a borrower.
    mp::owner<int*> owner{ptr};
    return owner;
}
```

mp::owner

Use `std::unique_ptr` or `std::shared_ptr` instead

mp::owner

- `as_borrower`
`borrower<T>`
`as_borrower() const {`
 `return borrower<T>(this->get());`
`}`
- Borrower explicit constructor
`template<Pointer U>`
`explicit borrower(owner<U> const &other)`

mp::borrower

- Compile time failure for delete
 - Can delete a pointer to a borrower
 - Can delete the pointer the borrower contains (BAD)
 - Cannot delete an object.
`mp::borrower<int*> p;`
`delete p;`

Backlog

- ~~Force check of maybe null~~
- What is the binary representation of 0

Nullptr binary representation

- `conv.ptr`
- `basic.compound`
- On most platforms, a null pointer is represented by a binary value of all zeros (e.g., `0x00000000` on a 32-bit system). However, this is not guaranteed by the standard. Some platforms might use a different representation for null pointers.

Backlog

- ~~What is the binary representation of 0~~

Summary int*

Three dimensions:

- Can be null?
- Collection?
- Ownership?

Summary Collection

Range
checking with
std::span



Summary int*

Nullptr is in-band
encoding



Summary int*

Ownership is not specified



imgflip.com

JANE-CLARK.TUMBLR

Summary int*

Wrappers give
strong type checking
at compile time



Thank you

marcpawl@gmail.com